

Part2- Section 3 (Deterministic vs Non Deterministic)

1. Based on the rule-based version, how easy was it to understand how equitably loan decisions are made using Themis?

- Very difficult
- Somewhat difficult
- Neutral
- Easy
- Very Easy

2. Based on the DL-based version, how easy was it to understand how equitably churn decisions are made using Themis?

- Very difficult
- Somewhat difficult
- Neutral
- Easy
- Very Easy

3. For the rule-based version, do Themis generated tests provide additional insights beyond what you can already see in the rules?

- Yes — it confirms discrimination assumptions and reveals corner-cases
- Maybe — rules give transparency, but testing quantifies bias and builds confidence.
- Only when rules are complex — simple rule systems may not benefit as much.
- Absolutely — it uncovers hidden patterns and supports explainability.
- No - It does not help

4. In the DL-based version, do you feel tests generated from Themis are sufficient to explain bias without seeing the internal model training structure?

- Yes — discrimination tests alone are sufficient to explain model bias.
- Mostly yes — they provide strong evidence, even without knowing internal training details.
- Somewhat — they reveal bias patterns, but additional model insights would be useful.
- Not really — they indicate bias but do not fully explain why the model behaves that way.
- No — understanding the internal training data and model design is necessary to explain bias.

5. Could Using Themis help you better understand how discrimination can occur in software outcomes?

- It did not help me understand discrimination at all
- It slightly helped, but I still found discrimination unclear
- It somewhat increased my understanding
- It helped me understand discrimination more clearly
- It significantly improved my understanding of how discrimination occurs in software outcomes

6. How much control did you feel you had over how Themis conducted and presented discrimination testing?

- The tool decided everything automatically; I couldn't influence the testing process
- I could make minor adjustments but had limited influence
- I could configure certain aspects but not the overall testing process
- I could meaningfully adjust how tests were conducted and interpreted
- I had full autonomy to modify and direct the testing and interpretation process

7. How useful do you find defining input subspaces (smaller, focused regions within a large input space) as a method for understanding and analyzing discrimination in complex models?

- Not useful at all – It provides little or no help in understanding discrimination
- Slightly useful – It helps in a limited way but lacks clarity or scalability
- Moderately useful – It provides some insight into discrimination patterns
- Useful – It meaningfully improves understanding of discrimination across inputs
- Extremely useful – It is a highly effective and insightful approach for analyzing discrimination in large input spaces

This content is neither created nor endorsed by Microsoft. The data you submit will be sent to the form owner.



Microsoft Forms